

Week 3

Spyros Mastorakis


Outline

- Primary backup replication
- Assignment 2

What to do with fail-stop failures?

- Last time: crash failures
 - Replicas can resume execution with saved state
- Today: fail-stop failures
 - All state is lost upon failure
 - Need to replicate state proactively
- Approach: primary backup replication
- Challenges
 - What if primary or backup fails?
 - How can we keep them in sync?

Primary Backup Replication



Key Idea: transparently have two replicas, primary and backup

- Primary interacts with client
- Backup stores copy of primary's state

Handling failures

- How to handle primary failure?
 - Promote one of the backups to be the primary
- How to handle backup failure?
 - Add another machine as a backup

When to sync?

- It is okay for the primary to be out of sync until a change is externally visible
 - External consistency: primary backup in sync to external world
- Implications for when primary should sync with backups?
 - Sync must happen before any state change is externally visible

What to transfer during sync


- Heavyweight: send snapshot of primary's state
 - **Slow!**
 - When is this necessary?
 - When bootstrapping a new backup
- Lightweight: send every operation
 - Why is this okay?
 - Leverage determinism of state machine
 - Send any state necessary for backup to mimic execution

Transparent Primary Backup

- Application relies on library to keep primary and backups in sync
- Library functions
 - Receive message from client
 - Sync with backups before sending response to client
- Will this solution work?
 - No! Does not capture nondeterminism in execution

VMM—based Primary Backup


- Primary and backup execute on two VMs
 - Primary logs inputs and outputs to log
 - Backup applies inputs from log
 - Primary waits for backup output
- Primary-backup monitor each other
 - If primary fails, backup takes over



Log-based VM replication

- Primary VMM sends log entries to backup VMM over the logging channel
- Backup VMM (hypervisor) replays log entries
 - Stops backup VM at next input event or nondeterministic instruction
 - Delivers same input event as primary used
 - Delivers same nondeterministic value as primary
- Ensures backup and primary state never diverge


Primary Backup Replication



Client perspective

- What does a worker need to know in order to register itself with the replicated master?
 - Needs to know which machine is primary
- Can the primary be hardcoded into client app code?
 - No! Primary will get replaced when it fails
- How does the client discover the current primary?
 - Needs reliable service to do primary lookups

Primary Backup Replication



View Service

- Maintains current membership of primary-backup service (i.e., **view**)
 - Each view → (view number, primary, backup)
- When does a view service change the view?
 - When primary or any backup fails
 - Periodically exchange heartbeat messages to detect failures
- What if view service is down or not reachable?

Transitioning between views

- View service broadcasts view change to all replicas
- Primary must ACK new view once backup is up to date
- Two implications
 - Liveness detection timeout > state transfer time
 - Cannot change view if primary fails during sync


View Service

- Summary: **view change has three steps**
 - View service announces new view to all replicas
 - Primary syncs with new backup if there is one
 - Primary acknowledges new view
- **Stuck if primary fails in the midst of this process**

Scalability of View Service

- Does every client need to contact view service before any operation?
 - No--clients can cache view across operations
- When to invalidate cached view?
 - Client invalidates cache when no response or negative response from replica it thinks is primary

Split Brain scenario



Avoiding Split Brain

- Primary must forward all operations to backups
 - Goal: get ACKs from backups that they too recognize primary
- Why can't backups be mistaken about who is primary?
 - Only a backup can be promoted as a primary

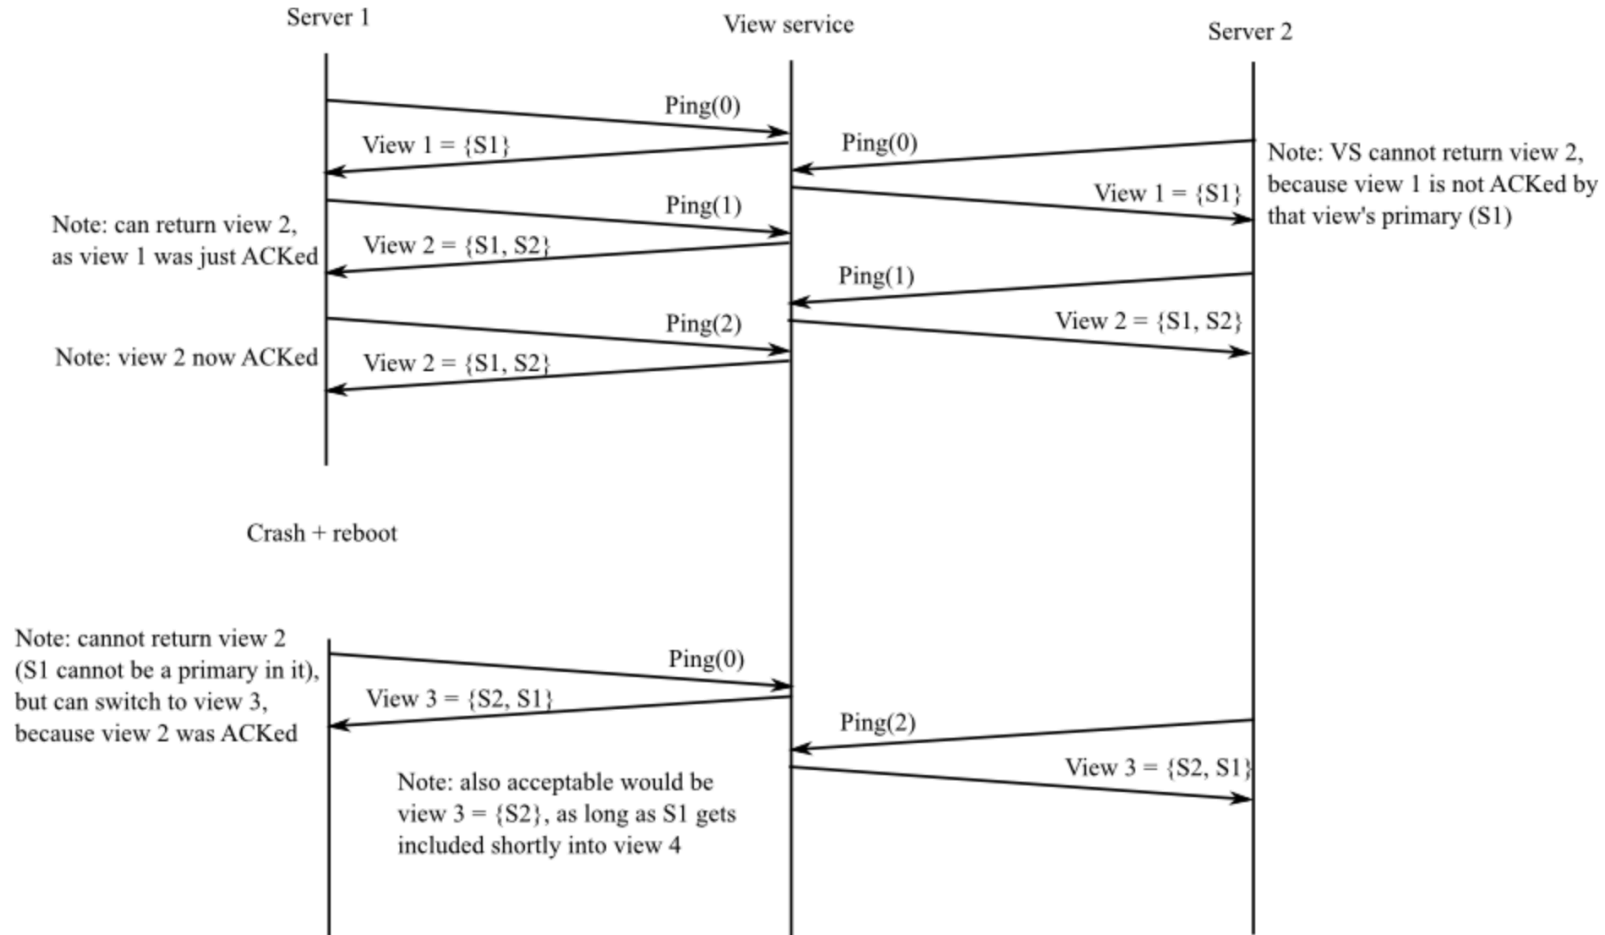
View Service Sequence Examples

- Valid sequence of views
 - $(1, S1, _) \rightarrow (2, S1, S2) \rightarrow (3, S1, S3) \rightarrow (4, S3, S4) \rightarrow (5, S4, _)$
- Examples of invalid transitions between views:
 - $(1, S1, S2) \rightarrow (2, S3, S4)$
 - $(1, S1, S2) \rightarrow (2, _, S2)$
 - $(1, S1, _) \rightarrow (2, S2, S1)$
 - $(2, S1, _) \rightarrow (1, S1, S2)$

Assignment 2

- Primary backup replication for a key/value service
- Coordination on who is the primary and who is a backup server through viewservice
- Viewservice monitors whether each server is alive or dead

Assignment 2 (cont'd)



Part A: Implementing the viewservice

- Based on previous figure, add field(s) to ViewServer in server.go in order to keep track of the most recent time at which the viewservice has heard a Ping from each server
- Viewservice keeps track of whether the primary for the current view has acknowledged it
- Viewservice needs to make periodic decisions, for example to promote the backup if the viewservice has missed DeadPings pings from the primary
- Study the test cases before you start programming
 - If you fail a test, you may have to look at the test code in test_test.go to figure out the failure scenario is

Part B: Primary backup key/value service

- **Assumption:** the viewservice never halts or crashes
- Implement the client and server parts
- Clients use the service by creating a Clerk object (see client.go) and calling its methods, which send RPCs to the service:
 - If the current primary does not respond, or doesn't think it's the primary, have the client consult the viewservice
- Server side:
 - Pings viewservice to find the current view
 - Primary forwards updates to backup
 - When a server becomes the backup in a new view, the primary should send it the primary's complete key/value database