


# Discussion 06

Spring 2019 – CS 188

Section 2B


# CAP Theorem

- Pick any two: consistency, availability, and partition-tolerance
  - In practice, choose between CP and AP



# PACELC

- If partition,
  - Choose availability vs. consistency
- Else,
  - Choose latency vs. consistency
- Unifies two separate tradeoffs that we've talked about



# Example1

- **Eventually consistent** key-value store.
  - When two replicas get near each other, they sync with each other.
  - Similar to Shared calendar application in the lecture.
  - If no new updates, all replicas **eventually** converge to same state.
- Format of updates?
  - $\langle (k1, v1), X \rangle$

# Updates

- Ordering of updates using Lamport clock
  - $\langle (k1, v1), X, 20 \rangle$  (20 is the timestamp)
- **Resolve conflicts** upon receiving updates from the other node
  - Node A:  $\langle (k1, v1), X, 20 \rangle$
  - Receives update from B:  $\langle (k1, v2), Y, 10 \rangle$
- To be able to “roll back”: maintain a log of updates in each replica.
  - Node A:  $\langle (k1, v0), X, 5 \rangle \langle (k2, v2), X, 15 \rangle \langle (k1, v1), X, 20 \rangle$
  - Re-sort the log again
  - Replay

# Remember: Quick Syncs

- How to sync without state exchange proportional to size of log?

A	B
$\langle -, 10, X \rangle$	$\langle -, 10, X \rangle$
$\langle -, 20, Y \rangle$	$\langle -, 20, Y \rangle$
$\langle -, 30, X \rangle$	$\langle -, 30, X \rangle$
$\langle -, 40, X \rangle$	$\langle -, 40, Y \rangle$

- B tells A: **highest timestamp for every node**
  - E.g., “X 30, Y 40”  $\leftarrow$  **Version Vector**
  - In response, A sends all of X’s updates after (, 30, X) and all of Y’s updates after (, 40, Y)

# Sync

- Sync updates, not the state!
- Arguments?
  - Version vector: [X:20, Y:10]
- Return Value?
  - List of updates the caller does not have
  - All updates from X after 20 and from Y after 10.