


Discussion 07

Spring 2019 – CS 188


Section 2B

Consistent Hashing

- Represent hash space as a circle
- Partition keys across servers
 - Assign every server a random ID
 - Hash server ID
 - Server responsible for keys between predecessor and itself
- How to map a key to a server?
 - Hash key and execute read/write at successor




Successor Pointers




- If you don't have value for key, forward to successor

Finger Tables


- i 'th entry at node n points to successor of $\text{hash}(n) + 2^i$
 - # of entries = # of bits in hash value
- Binary lookup tree rooted at every node
 - Threaded through others' finger tables



Example1: Lookup with Finger Table



Lookups take $O(\log N)$ hops




Example2: Serializability

- Concurrent execution of transactions:
 - T1: transfer \$10 from Alice to Bob
 - T2: Read balance in Alice's and Bob's accounts
 - T3: Read balance in Charlie's account.
 - Initial balance is \$100 of 3 accounts
- Permissible outputs for T2?
 - (Alice: \$100, Bob: \$100) or (Alice: \$90, Bob: \$110)
- Invalid outputs for T2?
 - (Alice: \$90, Bob: \$100) or (Alice: \$100, Bob: \$110)


Transactions and Serializability

- Concurrent execution of transactions:
 - T1: transfer \$10 from Alice to Bob
 - T2: Read balance in Alice's and Bob's accounts
 - T3: Read balance in Charlie's account.
 - Initial balance is \$100 of 3 accounts
- From user's perspective: these transactions output must be the same as some serial execution.

Example2: Serializability



Some of the possible Orderings:



Two Phase Locking

